

SEQR: searching protein sequences by similarity in Solr

Lewis Geer


National Institutes of Health



Questions

- Algorithm
 - Can we create a similarity search that is computationally efficient? Space efficient?
- Implementation
 - Can we improve the usability of sequence search in the era of nextgen sequencing?
 - Can we implement the algorithm by leveraging open source to minimize development costs?


Similarity indexes become necessary as sequencing beats Moore's law



BLAST 2006

Distribution of 33 Blast Hits on the Query Sequence


Mouse-over to show defline and scores, click to show alignments



BLAST 2015

Distribution of 101 Blast Hits on the Query Sequence

Mouse over to see the define, click to show alignments




How do we provide useful results quickly?



- Surface interesting sequences first
- Integrate queries into search results
 - Allow useful sub-setting after the search, e.g. taxonomy
 - Faceting
 - Autosuggestion
- Display information in more understandable ways, e.g. trees.
- Quick response to allow iterative discovery
- Reconfigure interface to fit specific user needs

Demo


Fly Wnt8 protein



Human Wnt8 protein




Fly embryo layout



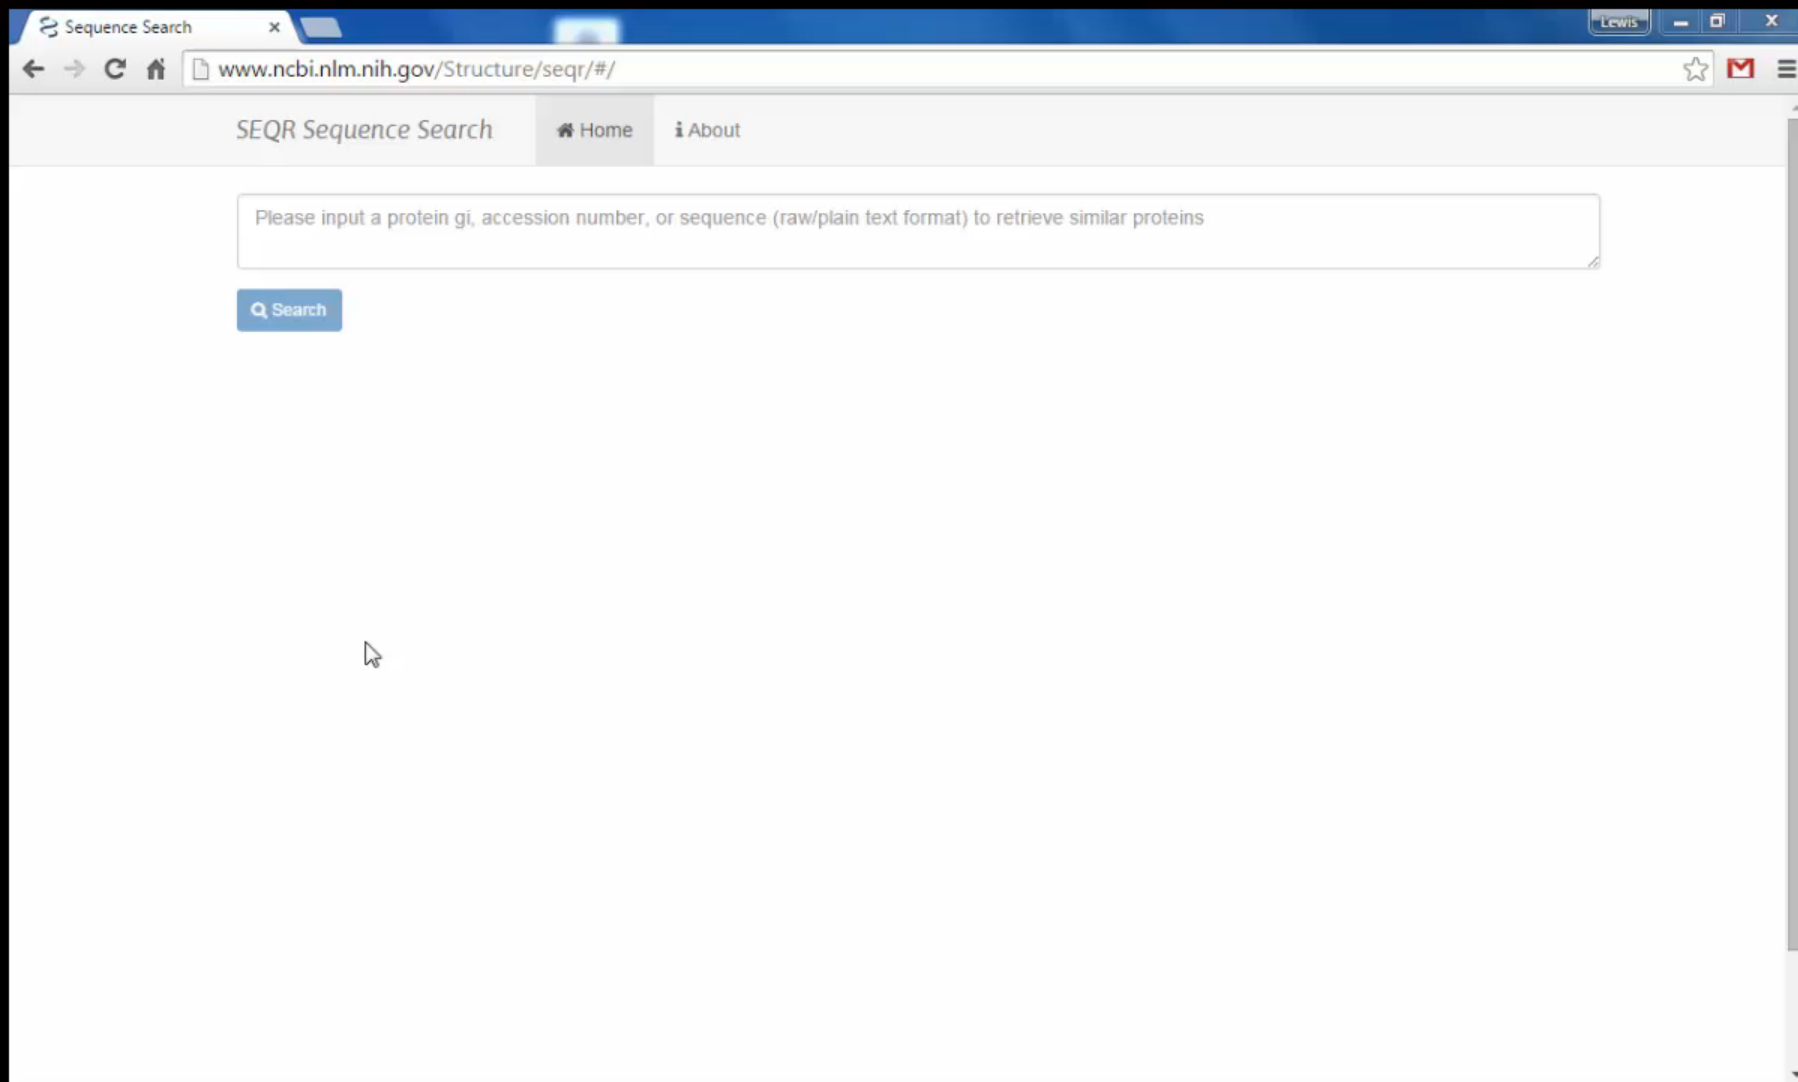
Human kidney
Grown in lab

Melissa Little et al.



Human	KAYLTYSASVAVGAQNGIEECKYQFAWERWNCPESTLQLATHNGLRSATRETSFVHAIS
Fly	+A L++ G + ++ C+ F W+RWNCPE + +Q + S RE +V AIS
	QAPLSWEDITGKGLKQALDSCQSFQWRWNCPSQDFVQKNSKPEENSPNREDVYVAIS
	SAGVMTLTRNCMGMDFDNCGDDSRNGRIGGRGWVWGGCSDNAEFGERISKLFVDGLET
	A +++TLT++C+ G CGC ++ + + E+ K F
	MAAIVHTLTKDCANGVIAGCGCTENALNVPCAH-----EPTKALEQYEKHF-----
	GQDARALMNLHNEAGRLAVKETMKRTCKCH---GISGSCSIQTCWLQLAEFRDIGNHLK
	G + A+ HN ++ ++++ C+C + G C + C L F I L
	GSGSGAIG--HNRRVVGALLQRSLEQECRCCKQPGAVQGECQEEECVAVLKPF E A I A Q D L L
	IKHDQALKLEMDKRKMRSNSADNRGAIADAFSSVAGSELIFLEDSPDYCLKNISLGLQG
	+D A++LE G S++ + + ++ L+F++DSP+YC ++ + GL G
	QMYDDAIQLE-----GASSN----LKIMWQNIPLDSLVMQDSPNYCERDAT-GLWG

<http://www.ncbi.nlm.nih.gov/Structure/seqr/>

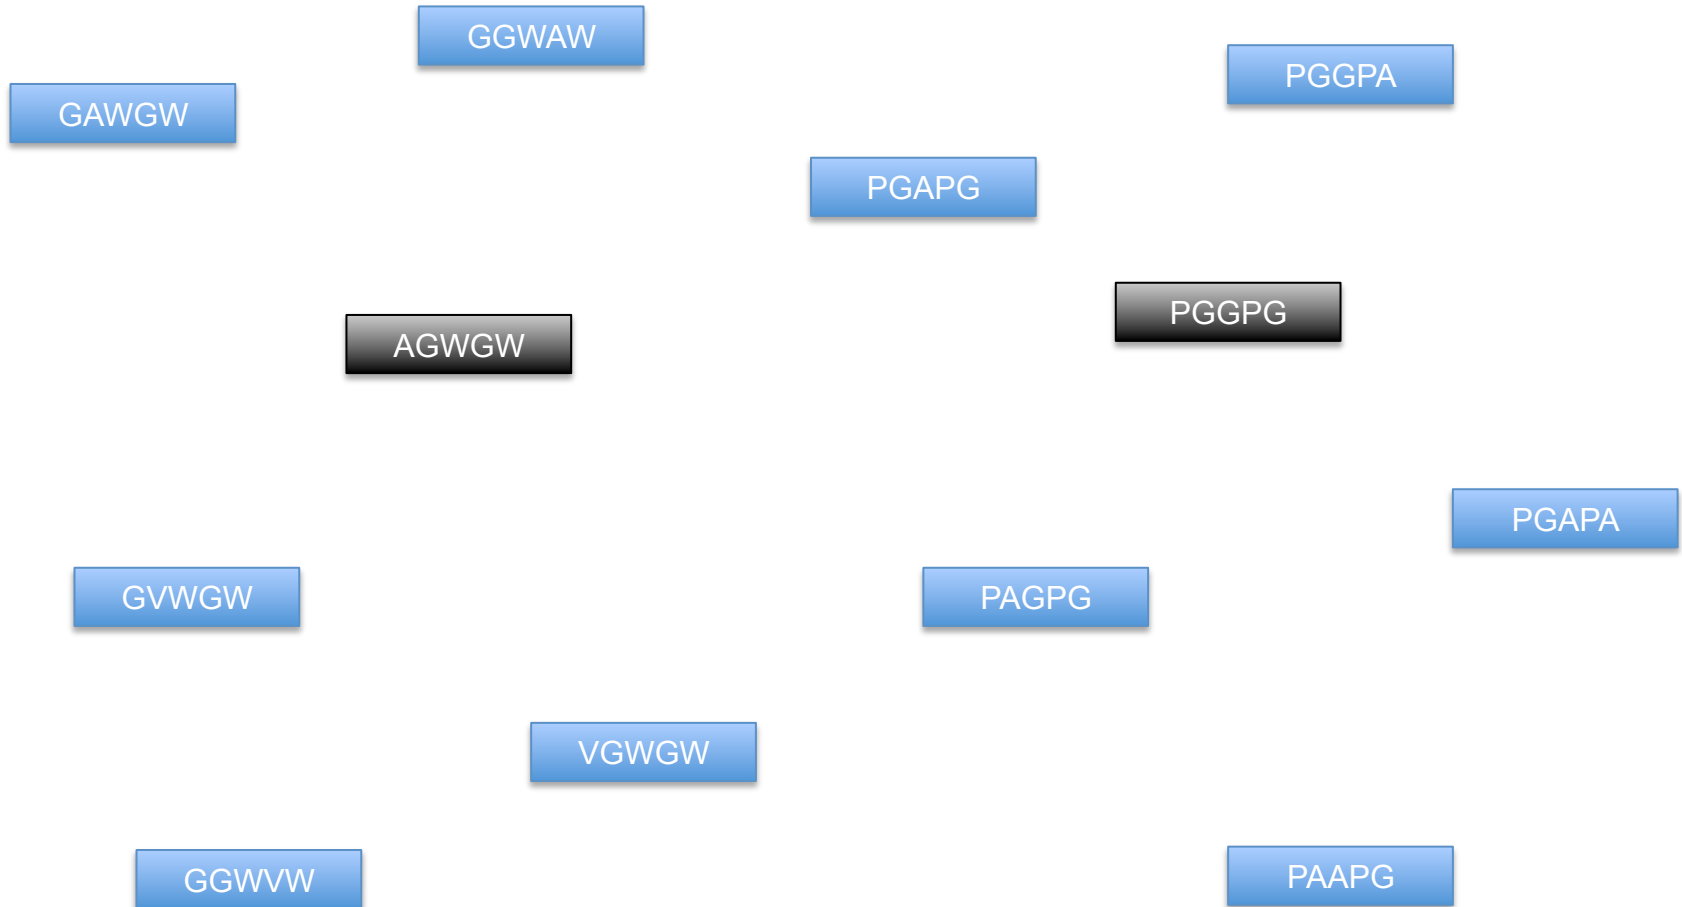


ALGORITHM

Generate all nmers



Select cluster seeds



Assign PSSM to seed nmers


Substitution matrix

	A	G	P	R	W
A	4	0	-1	-1	-3
G	0	6	-2	-2	-2
P	-1	-2	7	-2	-4
R	-1	-2	-2	5	-3
W	-3	-2	-4	-3	11

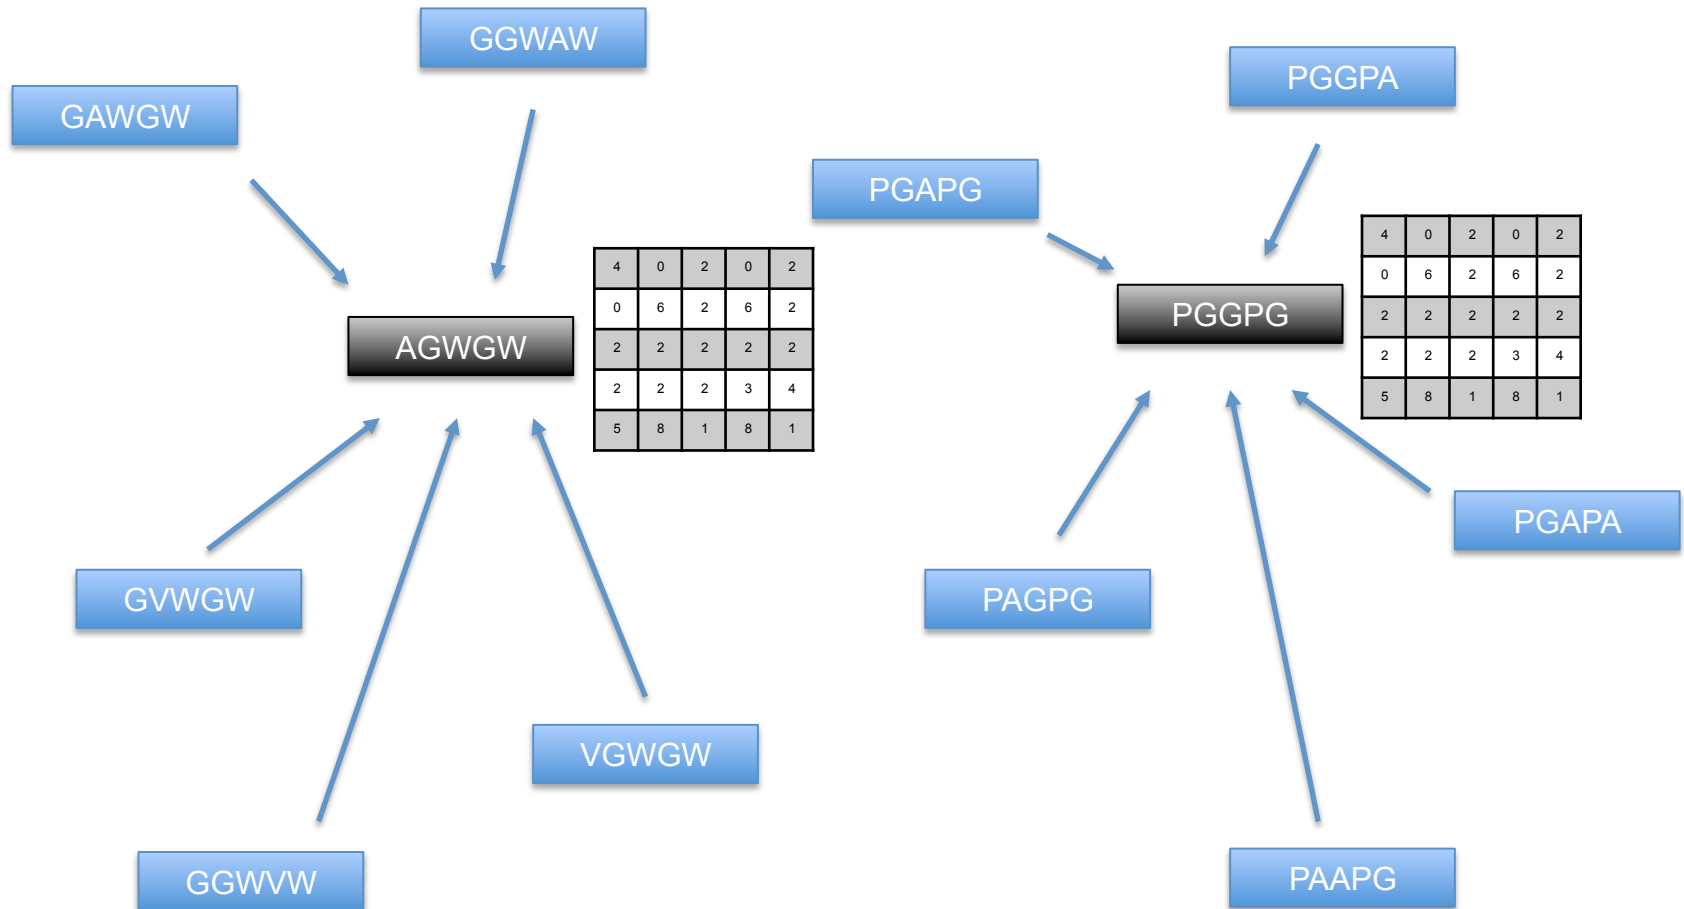
Seed PSSM

4	0	-3	0	-3
0	6	-2	6	-2
-1	-2	-4	-2	-4
-1	-2	-3	-2	-3
-3	-2	11	-2	11

AGWGW



Cluster all other nmers based on distance to seed PSSM w/ score threshold



Further steps in clustering

- Seed PSSMs can be updated to average in the PSSMs of merged members
- After complete assignment, can reassign members based on averaged pssms
- Assign each cluster an ID.

Sequence indexing

- Create an nmer to cluster ID index. Sequences can be indexed in microseconds. This is an optional step.
- Scan each database sequence, position by position. Assign cluster ID to each position either by using index or direct comparison to PSSM.


Retrieval

- Create inverted index from cluster ID to sequence ID
- Retrieval procedure:
 - Database is preindexed
 - Query is indexed on the fly
 - For each cluster ID from query, retrieve the sequence IDs and keep track of which sequence IDs have hits

Retrieval, continued...

- After processing all cluster IDs in query, rank hit sequences by Jaccard score: $|A \cap B| / (|A| + |B| - |A \cap B|)$
- Note that this is a retrieval algorithm, not an alignment algorithm. Alignment can be applied as a post processing step.

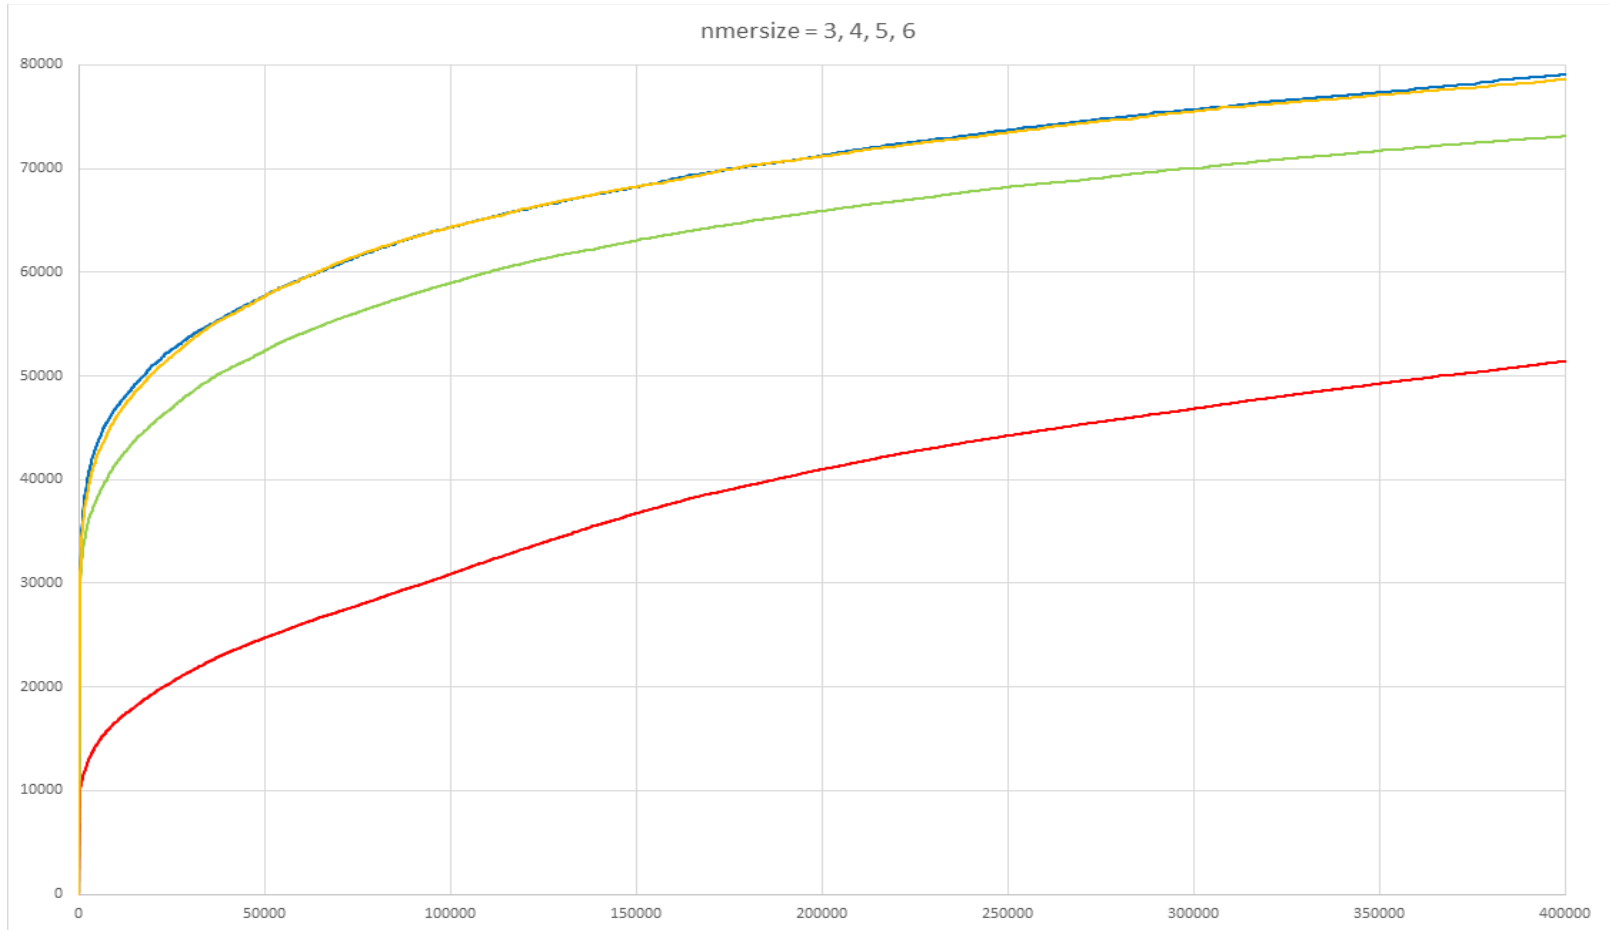
Processing speed



ROC_n analysis


- Compare ~300 query sequences selected from Swiss-Prot, with a bias toward human sequences
- Search against Swiss-Prot (450k sequences)
- Compare to DELTA-BLAST with a threshold cutoff of $1e-4$.

Size of nmer




red: n-mer size = 3, threshold = 5.2
green: n-mer size = 4, threshold = 4.6
blue: n-mer size = 5, threshold = 3.6
orange: n-mer size = 6, threshold = 2.8

Vary threshold



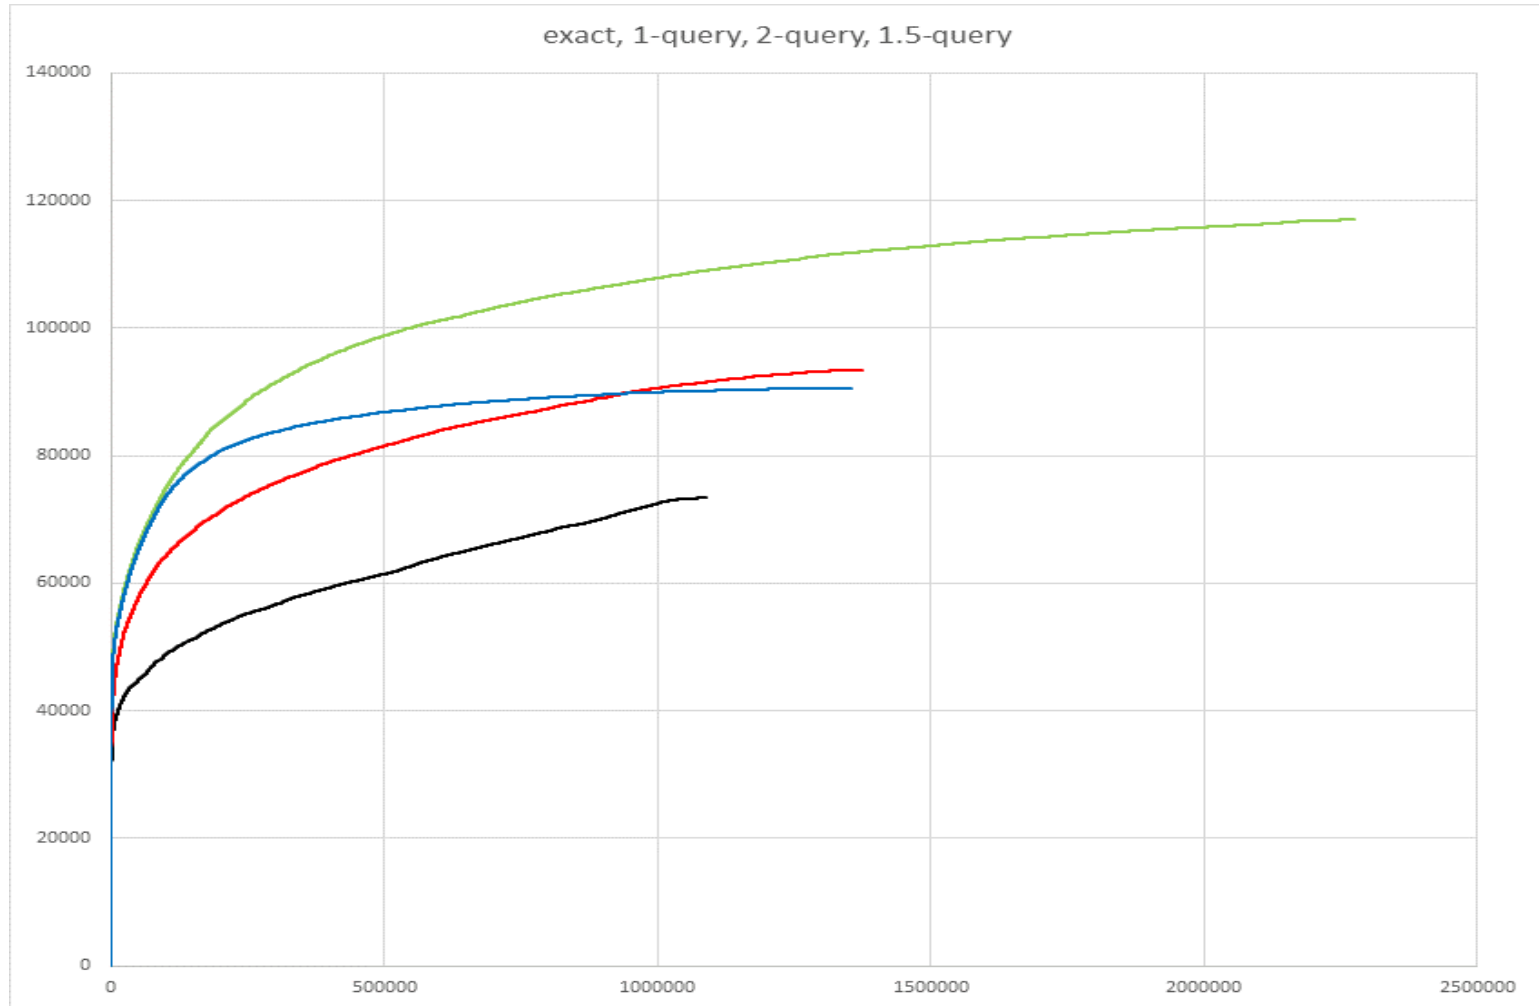
red: 3.5, green: 3.6, blue: 3.7, orange: 3.8, magenta: 3.9
red-dashed: 4.0, green-dashed: 4.1, black: exact

Substitution matrix



red: blosum45, green: blosum50, blue: blosum62, orange: blosum80, magenta: blosum90
red-dashed: pam30, green-dashed: pam70, blue-dashed: pam250

Number of queries



These ROC curves show the improvement of optimized 2-query SEQR (green) vs our best 1-query SEQR (red) and exact-match indices (black)

Cluster sizing

- Best performing indices have ~150k clusters with an average size of ~20

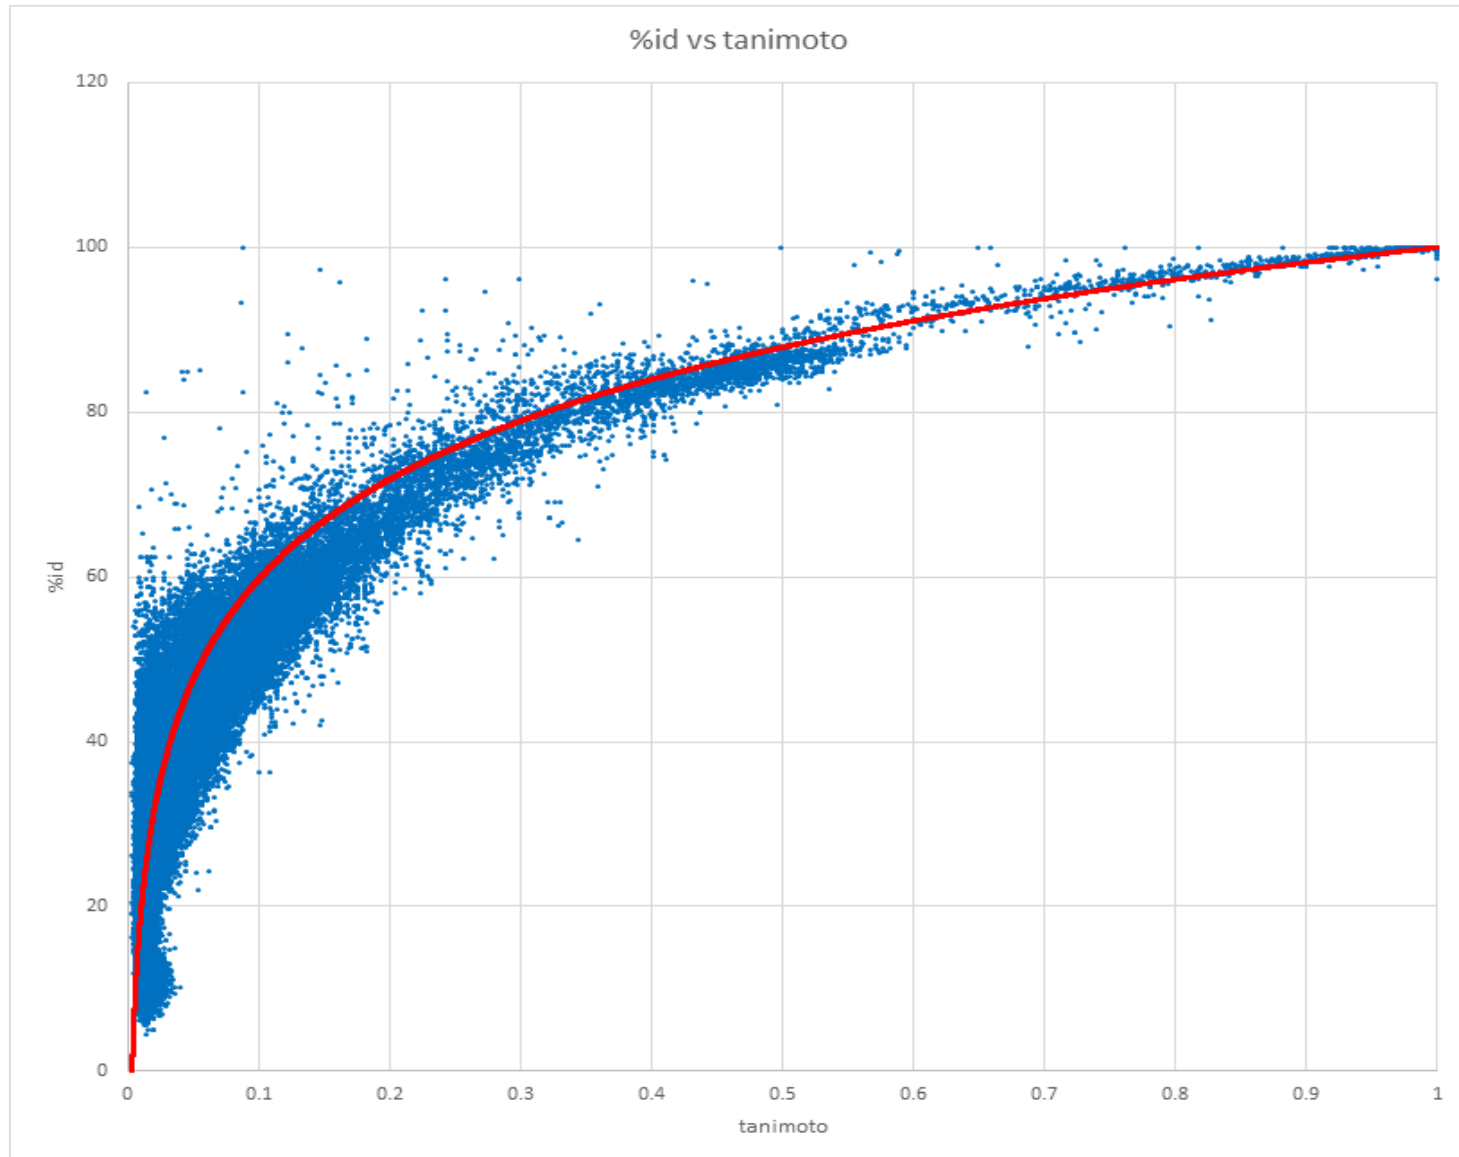
0 - 0 : 5	60 - 69 : 2149
1 - 1 : 1314	70 - 79 : 1274
2 - 2 : 4433	80 - 89 : 984
3 - 3 : 5785	90 - 99 : 815
4 - 4 : 13265	100 - 199 : 3131
5 - 5 : 6814	200 - 299 : 924
6 - 6 : 10643	300 - 399 : 367
7 - 7 : 7268	400 - 499 : 72
8 - 8 : 14519	500 - 599 : 47
9 - 9 : 5922	600 - 699 : 25
10 - 19 : 42642	700 - 799 : 16
20 - 29 : 14647	800 - 899 : 11
30 - 39 : 7724	900 - 999 : 11
40 - 49 : 4336	1000 - 1999 : 11
50 - 59 : 2407	

Sample cluster, seed FFTLT

FFTLT
FFTLI
FFTVT
FFVLT
FFKLT
YFTMT
FFTYT
FFRLT
FFSLT
FFTMT
FFTLQ
FFTLM
FYTLT
FFTLS

FFTLN
FFTLP
YFTLT
FFALT
FFTAT
FFTFT
FYTMT
FYTIT
FFTCT
FFTLV
FFTLL
FFELT
FFTTT
FFTIT

Compare BLAST hits to SEQR



Example hit

```
Query 185 LQRHRIMHRGDGPYKCKFCGKALMFLSLYLIHKRTHTEK-----P 225
          L  +++  +GD      K C K      + +      +  TGEK                      P
Sbjct 195 LPNNKLSDKGDKNQTSKKCEKVCRRHSASHTKEDKIQTGEKRRKSHCRTPSKPEKAPGSGKP 254

Query 226 YQCKQCGKAFSSSSSLRIHERHTHTEKPYKCNECGKAFHSSSTCLHAHKRTHTEKPYECK 285
          Y+C  CGK  SH   L  H+RTHTEKPY+CNECG AF   + L  H+RTHTEKPYEC+
Sbjct 255 YECNHCGKVLSHKQGLLDHQRTHTGEKPYECNECGIAFSQKSHLVVHQRTHTGEKPYECE 314

Query 286 QCGKAFSSSHSFQIHERHTHTEKPYECKECGKAFKCPSSVRRHERTHSRKKPYECKHCGK 345
          QCGKA      H+   H R HTGEKPY+C ECGK F+  S++ +H R+H+  +KPYECK CGK
Sbjct 315 QCGKAHGHHKHALTDHLRIHTGEKPYKCNECGKTFRHSSNLMQHLRSHTGEKPYECKECGK 374

Query 346 VLSYLTSTFQNHLMHTGEI SHKCKICGKAFYSPSSLQTHEKTHTEKPYKCNQCGKAFNS 405
          Y +S   H+   HTGEI ++C  CGKAF   SSL  H +  HTGEKP++CN+CGK F+
Sbjct 375 SFRYNSSLTEHVRTHTEGEI PYECNECGKAFKYGSSSLTKHMRHTGEKPFECNECGKTFSK 434

Query 406 SSSFRYHERHTHTEKPYECKQCGKAFRSASLLQTHGRTHTEKPYACEKCGKPFNSNFSFF 465
          S      H+RTHT EKPY+C +CGKAF  +S L   H RHTG+ P+ C +CGK F
Sbjct 435 KSHLVIHQRTHTKEKPYKCDECGKAFGHSSSLTYHMRTHTGDCPFECNQCGKAFKQIEGL 494

Query 466 QIHERMHREEKPYECKGYGKTFSLPSLFHRHERHTTGKTYECKQCGRSFNCSSSFRYHG 525
          H+R+H  EKPYEC   GK FS  S      H+RHTG K +EC +CG++FN  S      H
Sbjct 495 TQHQRVHTGEKPYECVECGKAFSQKSHLIVHQRTHTGEKPFECYECGKAFAKSQLVHIHQ 554

Query 526 RTHTEKPYECKQCGKAFRSASQLQIHGRTHTEKPYE 563
          R+HTGEKPYEC +CGKAF+  + L  H + H+ E+  E
Sbjct 555 RSHTGEKPYECIECGKAFKQNASLTKHMKIHSEEQSEE 592
```

- BLASTP percent identity: 37%
- SEQR tanimoto: 0.2
- 213 SEQR indexing terms
- For identical n-mers, 34 indexing terms

Typical hit found by SEQR but not by blastp

BLAST® Basic Local Alignment Search Tool

Home Recent Results Saved Strategies Help

My NCBI [Sign In] [Register]

NCBI/BLAST/blastp suite-2sequences/Formatting Results - FM3G4RP6114

[Edit and Resubmit](#) [Save Search Strategies](#) [Formatting options](#) [Download](#) [YouTube](#) [How to read this page](#) [Blast report description](#)

Blast 2 sequences

XP_005268682:PREDICTED: arf-GAP with GTPase,...

RID	FM3G4RP6114 (Expires on 03-08 04:28 am)	Subject ID	gi 4506509 ref NP_002917.1
Query ID	gi 530399716 ref XP_005268682.1	Description	regulator of G-protein signaling 12 isoform 2 [Homo sapiens]
Description	PREDICTED: arf-GAP with GTPase, ANK repeat and PH domain-containing protein 2 isoform X2 [Homo sapiens]	Molecule type	amino acid
Molecule type	amino acid	Subject Length	1376
Query Length	1172	Program	BLASTP 2.2.31+ Citation

Other reports: [Search Summary](#) [Taxonomy reports](#) [Multiple alignment](#)

☐ **Graphic Summary**

Distribution of 5 Blast Hits on the Query Sequence

Mouse-over to show define and scores, click to show alignments

Color key for alignment scores

<40	40-50	50-80	80-200	>=200
-----	-------	-------	--------	-------

Query

1 200 400 600 800 1000

IMPLEMENTATION

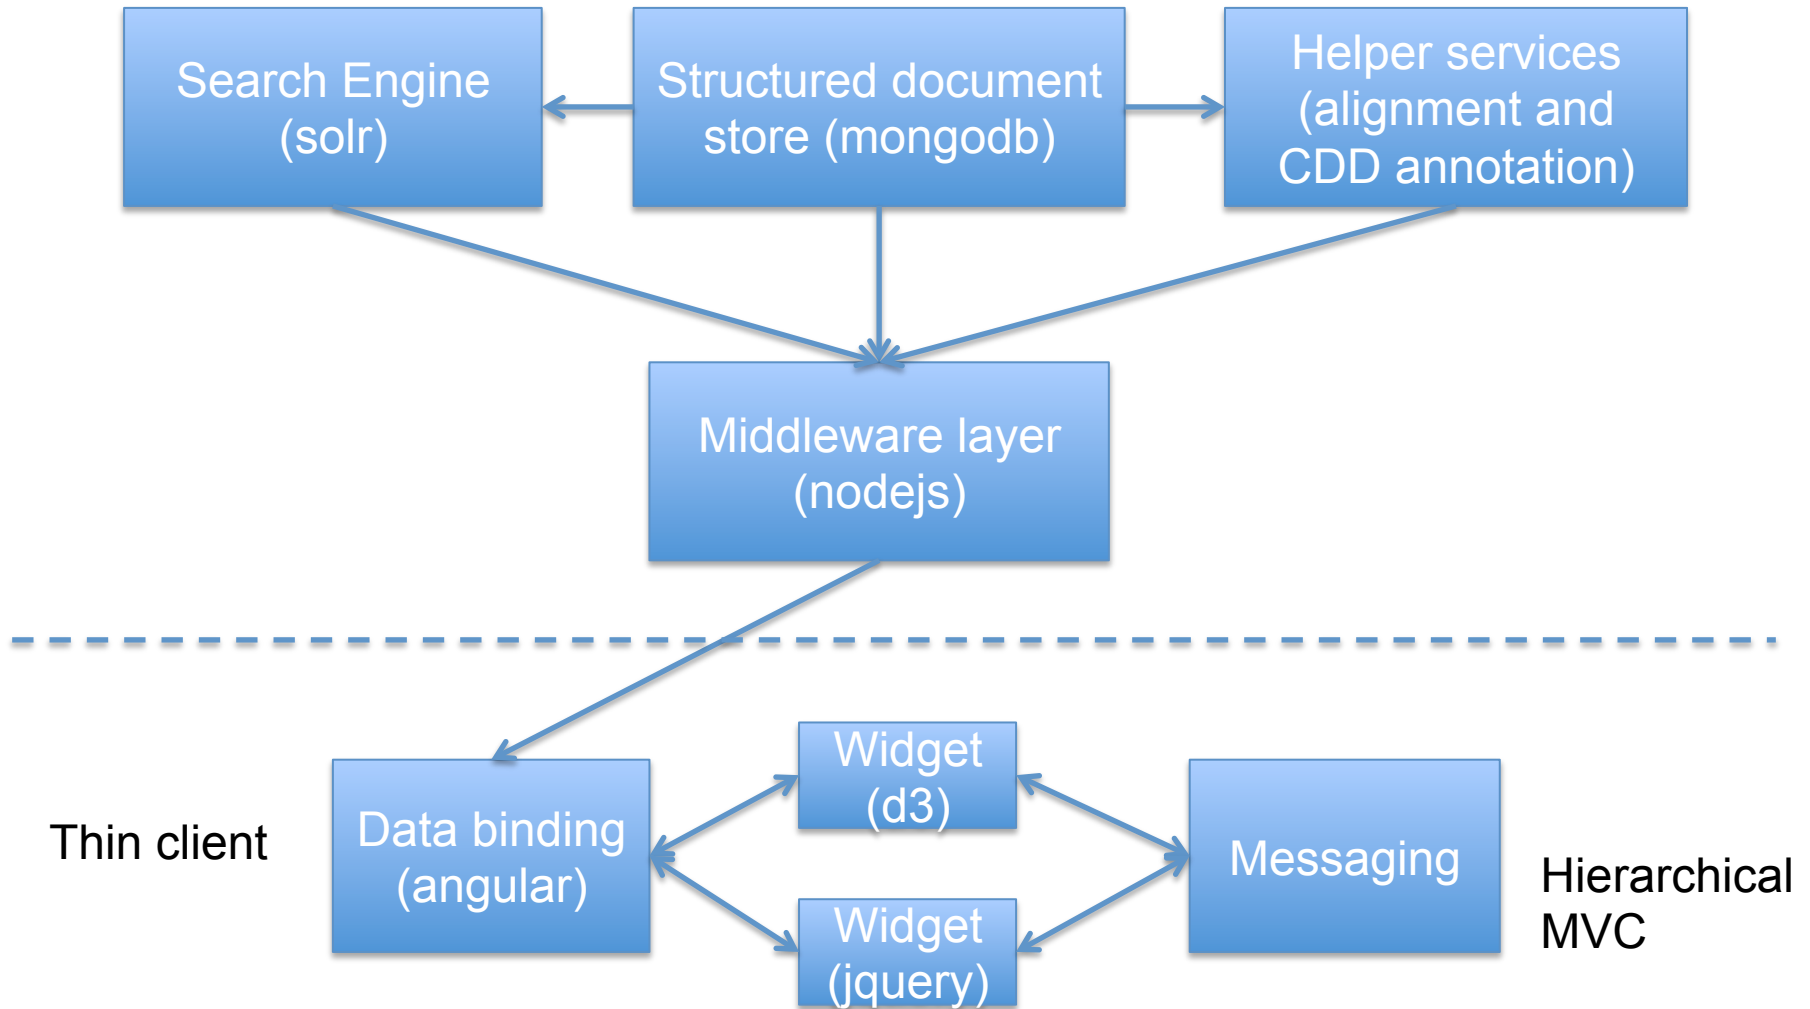
Implementation

- In the past few years, a quiet revolution in open source projects for information retrieval:
 - Search engines: solr, elasticsearch, ...
 - Structured data stores using JSON: mongodb, couchbase, ...
 - Middleware architectures with extensive libraries: nodejs, django, ...
 - Javascript: jquery, d3, angular
 - Even c++ has benefited: boost libraries, cereal, ...

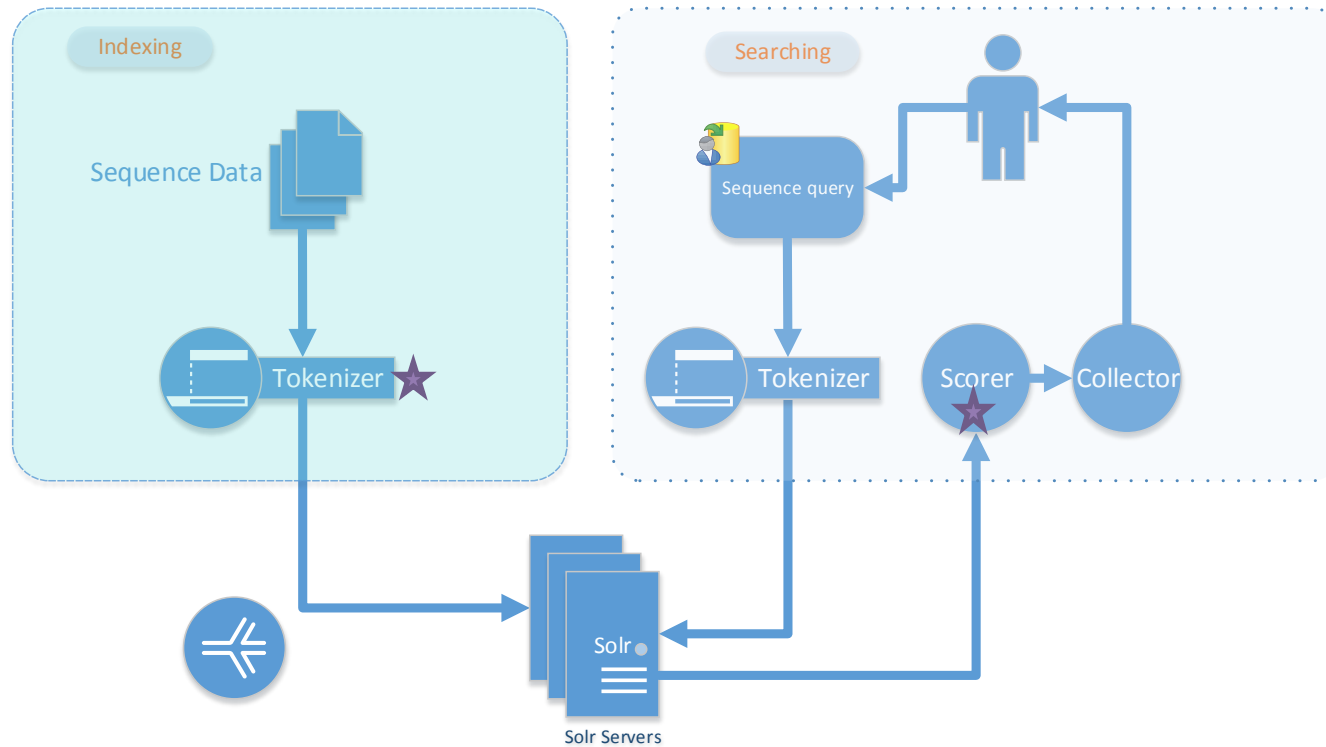
Implementation

- Web 2.0. Clean implementation of thin clients. The backend is SaaS.
- Can we use these libraries and techniques to speed up the development process (or make it possible for small groups to do big things)?

Architecture



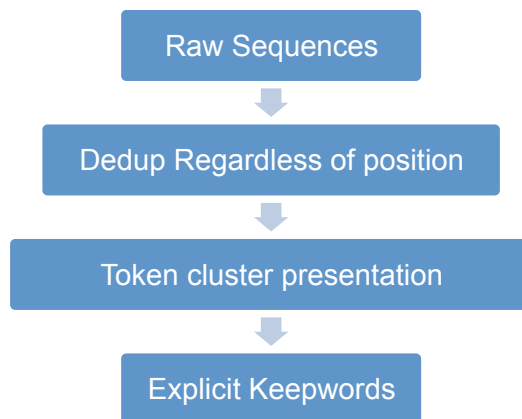
Workflow





Implementation

```
<fieldType name="sequence" class="solr.TextField" >
  <analyzer>
    <tokenizer class="gov.nih.nlm.ncbi.seqr.tokenizer.RawSequence2TokenizerFactory" nmer="5"/>
    <filter class="gov.nih.nlm.ncbi.seqr.tokenizer.RemoveDuplicatesTokenFilterIgnorePositionFactory"/>
    <filter class="solr.SynonymFilterFactory" synonyms="seqr.txt" ignoreCase="true" expand="false"/>
    <filter class="solr.KeepWordFilterFactory" words="seqr.keeplist" ignoreCase="true"/>
  </analyzer>
</fieldType>
```



<https://github.com/NCBI-Hackathons/seqr-tokenizer>

Tanimoto/Jaccard Index

```
<fieldType name="sequence" class="solr.TextField" >
  <similarity class="org.apache.lucene.analysis.tanimoto.OverlapSimilarity"/>
  <analyzer>
    <tokenizer class="gov.nih.nlm.ncbi.seqr.tokenizer.RawSequence2TokenizerFactory" nmer="5"/>
    <filter class="gov.nih.nlm.ncbi.seqr.tokenizer.RemoveDuplicatesTokenFilterIgnorePositionFactory"/>
    <filter class="solr.SynonymFilterFactory" synonyms="seqr.txt" ignoreCase="true" expand="false"/>
    <filter class="solr.KeepWordFilterFactory" words="seqr.keeplist" ignoreCase="true"/>
  </analyzer>
</fieldType>
```

```
q={!tanimoto bf=seqLen v=$qq}
qq={!edismax mm=50% qf=sequence}(bit1 bit2)
```

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|}.$$

<https://github.com/NCBI-Hackathons/solr-tanimoto>

Hardware

- Search running on one 2 CPU Dell C6220, 256 GB RAM, 1 TB SSD – low end commodity server
- $\frac{1}{4}$ of a 2U chassis.
- >200M sequences indexed. ~100 GB index size (< 1k/sequence)
- 15 cores used
- Typical 60 GB heap usage

Loading time

- Batch update of 236M records
 - Load MongoDB: 10.5 hours (1 thread).
Limited by use of one SQL-Server.
 - Dump MongoDB: 5 hours (20 threads).
Limited by number of MongoDB servers being queried (1).
 - Load Solr: 2.5 hours (16 threads).
 - Total time is 18 hours.

Acknowledgements

Lianyi Han

Shennan Lu

Jane He

Steve Bryant

Stephen Altschul

David Hurwitz

Bo Yu

Aron Marchler-Bauer

David Lipman

Future work

- Further analysis of algorithm
 - Clustering improvements
- Use of cluster pssms derived from CDD
 - Extension to domain annotation
- Extension to DNA sequences
- Alignment generation


```
>gi|665821313|ref|NP_001287868.1| protein Wnt-8a isoform 2 precursor [Homo sapiens]
MLCCIQCLCLVSPFPTLTPCQGGPHCLIPHLCLTFSLFGRSVNNFLITGPKAYLTYTTSVALGAQSGIE
ECKFQFAWERWNCPENALQLSTHNRLRSATRETSFIHAISSAGVMYITKNCSMGDFENCGCDGSNNGKT
GGHGWIIWGGCSDNVEFGERISKLFVDSLEKGDARALMNLHNNRAGRLAVRATMKRTCKCHGISGSCSIQ
TCWLQIAEFEREMGDYIKAKYDQALKIEMDKROLRAGNSAFEGHWWPAEAEIIPSAEAEIIEFESPDYCTCN
```

Search

Start Over











AAH26246: A2M proteinOrganism : [Homo sapiens](#)**Domain Classification:** This protein contains one functional domain "WNT1: found in Wnt-1". [More...](#)

All

[Protein w/ PubMed Reference](#)[3D Structure](#)[RefSeq](#)[Swiss-Prot](#)[Bioassay Target](#)[Phylogenetic tree](#)

Filters

Download

Identifier	Description	Organism	Alignment	Length	Identity	Actions
NP_001287868	protein Wnt-8a isoform 2 precursor	Homo sapiens		369	100%	Fasta , Seqr , Links
XP_008952854	PREDICTED: protein Wnt-8a isoform X1	Pan paniscus		369	99%	Fasta , Seqr , Links
XP_009448001	PREDICTED: protein Wnt-8a isoform X2	Pan troglodytes		369	99%	Fasta , Seqr , Links
NP_001287867	protein Wnt-8a isoform 1 precursor	Homo sapiens		386	96%	Fasta , Seqr , Links
XP_008952856	PREDICTED: protein Wnt-8a isoform X3	Pan paniscus		386	96%	Fasta , Seqr , Links
BAB60960	WNT8A	Homo sapiens		351	96%	Fasta , Seqr , Links
NP_490645	protein Wnt-8a isoform 3 precursor	Homo sapiens		351	96%	Fasta , Seqr , Links
EAW62166	wingless-type MMTV integration site family, member 8A, isoform CRA_b	Homo sapiens		351	96%	Fasta , Seqr , Links
AAI56845	Wingless-type MMTV integration site family, member 8A	synthetic construct		351	96%	Fasta , Seqr , Links
Q9H1J5	Protein Wnt-8a	Homo sapiens		351	96%	Fasta , Seqr , Links

1

Page 1 of 998

10

View 1 - 10 of 9,974

```
>gi|665821313|ref|NP_001287868.1| protein Wnt-8a isoform 2 precursor [Homo sapiens]
MLCCIQCLCLVSPFPTLTPCQGGPHCLPIHLCLTFSLFGRSVNNFLITGPKAYLTYTTSVALGAQSGIE
ECKFQFAWERWNCPENALQLSTHNRLRSATRETSFIHAISSAGVMYIITKNCMGMDFENCDCGDSNNGKT
GGHGWIWGGCSDNVEFGERISKLFVDSLEKGDARALMNLHNNRAGRLAVRATMKRTCKCHGISGSCSIQ
TCWLQLAFFEREMGDYIKAKYDOALKIEMDKROLRAGNSAEFGHWPAFAELPSAEAELEIESPDYCTCN
```

🔍 Search

↺ Start Over

AAH26246: A2M proteinOrganism : [Homo sapiens](#)**Domain Classification:** This protein contains one functional domain "WNT1: found in Wnt-1" . [More...](#)

All

Protein w/ PubMed Reference

3D Structure

RefSeq



Swiss-Prot

Bioassay Target

Phylogenetic tree

▼ Filters

📄 Download

	Protein	Alignment	Length	Identity	Actions
—	NP_001287868 : protein Wnt-8a isoform 2 precursor [Homo sapiens]		369AA	100%	Fasta , Seqr , Links
	Wnt family proteins are secreted and associated with the cell surface. Molecular cloning and characterization of human WNT8A. Expression and regulation of WNT8A and WNT8B mRNAs in human tumor cell lines: up-regulation of WNT8B mRNA by beta-estradiol in MCF-7 cells, and down-regulation of WNT8A and WNT8B mRNAs by retinoic acid in NT2 cells. Link to NCBI PubMed More...				
—	NP_001287867 : protein Wnt-8a isoform 1 precursor [Homo sapiens]		386AA	96%	Fasta , Seqr , Links
	Wnt family proteins are secreted and associated with the cell surface. Molecular cloning and characterization of human WNT8A. Expression and regulation of WNT8A and WNT8B mRNAs in human tumor cell lines: up-regulation of WNT8B mRNA by beta-estradiol in MCF-7 cells,				

```
>gi|665821313|ref|NP_001287868.1| protein Wnt-8a isoform 2 precursor [Homo sapiens]
MLCCIQCLCLVSPFPTLTPCQGGPHCLIPHLCLTFSLFGRSVNNFLITGPKAYLTYTTSVALGAQSGIE
ECKFQFAWERWNCPENALQLSTHNRLRSATRETSFIHAISSAGVMYIITKNCSMGDFENCGCDGSNNGKT
GGHGWIWGGCSDNVEFGERISKLFVDSLEKGDARALMNLHNNRAGRLAVRATMKRTCKCHGISGSCSIQ
TCWLQLAEFERMGDYIKAKYDOALKIEMDKROLRAGNSAEFGHWPAFAELPSAEAELEIEESPDYCTN
```

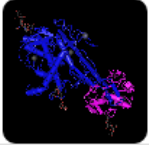

🔍 Search

↺ Start Over

AAH26246: A2M proteinOrganism : [Homo sapiens](#)**Domain Classification:** This protein contains one functional domain "WNT1: found in Wnt-1" . [More...](#)[All](#)[Protein w/ PubMed Reference](#)[3D Structure](#)[RefSeq](#)[Swiss-Prot](#)[Bioassay Target](#)[Phylogenetic tree](#)

▼ Filters

📄 Download

Structure	Alignment	Length	Identity	Actions
 <div>4F0AB - Xenopus laevis Chain B, Crystal Structure Of Xwnt8 In Complex With The Cysteine-rich Domain Of Frizzled 8</div>		316AA	70%	Fasta , Seqr , Links


🔍

⏪ ⏩ Page 1 of 1 ⏪ 10 ▼

View 1 - 1 of 1

3D Structure : 390981211

Tools



OK

download

Links

1 - 1 of 1

Sequence Search

←

→

↺

🏠

LQNSHNTSRWERRSCGRLCTECGLQVEERKTEVISSCNCKFQWCCTVKCDQCRHVVS%0AKYYCARSPGSAQSLGKGSA/all

☆

✉

☰

SEQR Sequence Search

🏠 Home

ℹ About

⬆ Filters 1 ✕

Tags

☐ Annotated (20)
☐ Gene (30)
☐ Gene Representative (6)
☐ NR Representative (17)
☐ OMIM (1)
☐ PubMed (10)
☐ Reference (10)

Source

☐ EMBL (2)
☐ GenBank (24)
☐ RefSeq (10)
☐ Swiss-Prot (4)

Organism

drosophila melanogaster ✕

Description


filter by description


Gene Symbol


filter by gene symbol

⬇ Download

Identifier	Description	Organism	Alignment	Length	Identity	Actions
AHN56049	Wnt oncogene analog 2, isoform B	Drosophila melanogaster	<div><div></div></div>	325	42%	Fasta , Seqr , Links
NP_001286251	Wnt oncogene analog 2, isoform B	Drosophila melanogaster	<div><div></div></div>	325	42%	Fasta , Seqr , Links
CAA46001	Wnt-2 protein	Drosophila melanogaster	<div><div></div></div>	352	42%	Fasta , Seqr , Links
AAO24959	RE36604p	Drosophila melanogaster	<div><div></div></div>	352	42%	Fasta , Seqr , Domains
AAF58933	Wnt oncogene analog 2, isoform A	Drosophila melanogaster	<div><div></div></div>	352	42%	Fasta , Seqr , Links
NP_476810	Wnt oncogene analog 2, isoform A	Drosophila melanogaster	<div><div></div></div>	352	42%	Fasta , Seqr , Links
P28465	Protein Wnt-2	Drosophila melanogaster	<div><div></div></div>	352	42%	Fasta , Seqr , Links
AFH97153	FI20276p1	Drosophila melanogaster	<div><div></div></div>	352	42%	Fasta , Seqr , Links
AAV55697	IP02562p	Drosophila melanogaster	<div><div></div></div>	304	36%	Fasta , Seqr , Domains







Comparison to delta BLAST

