



building better search for bioinformatics

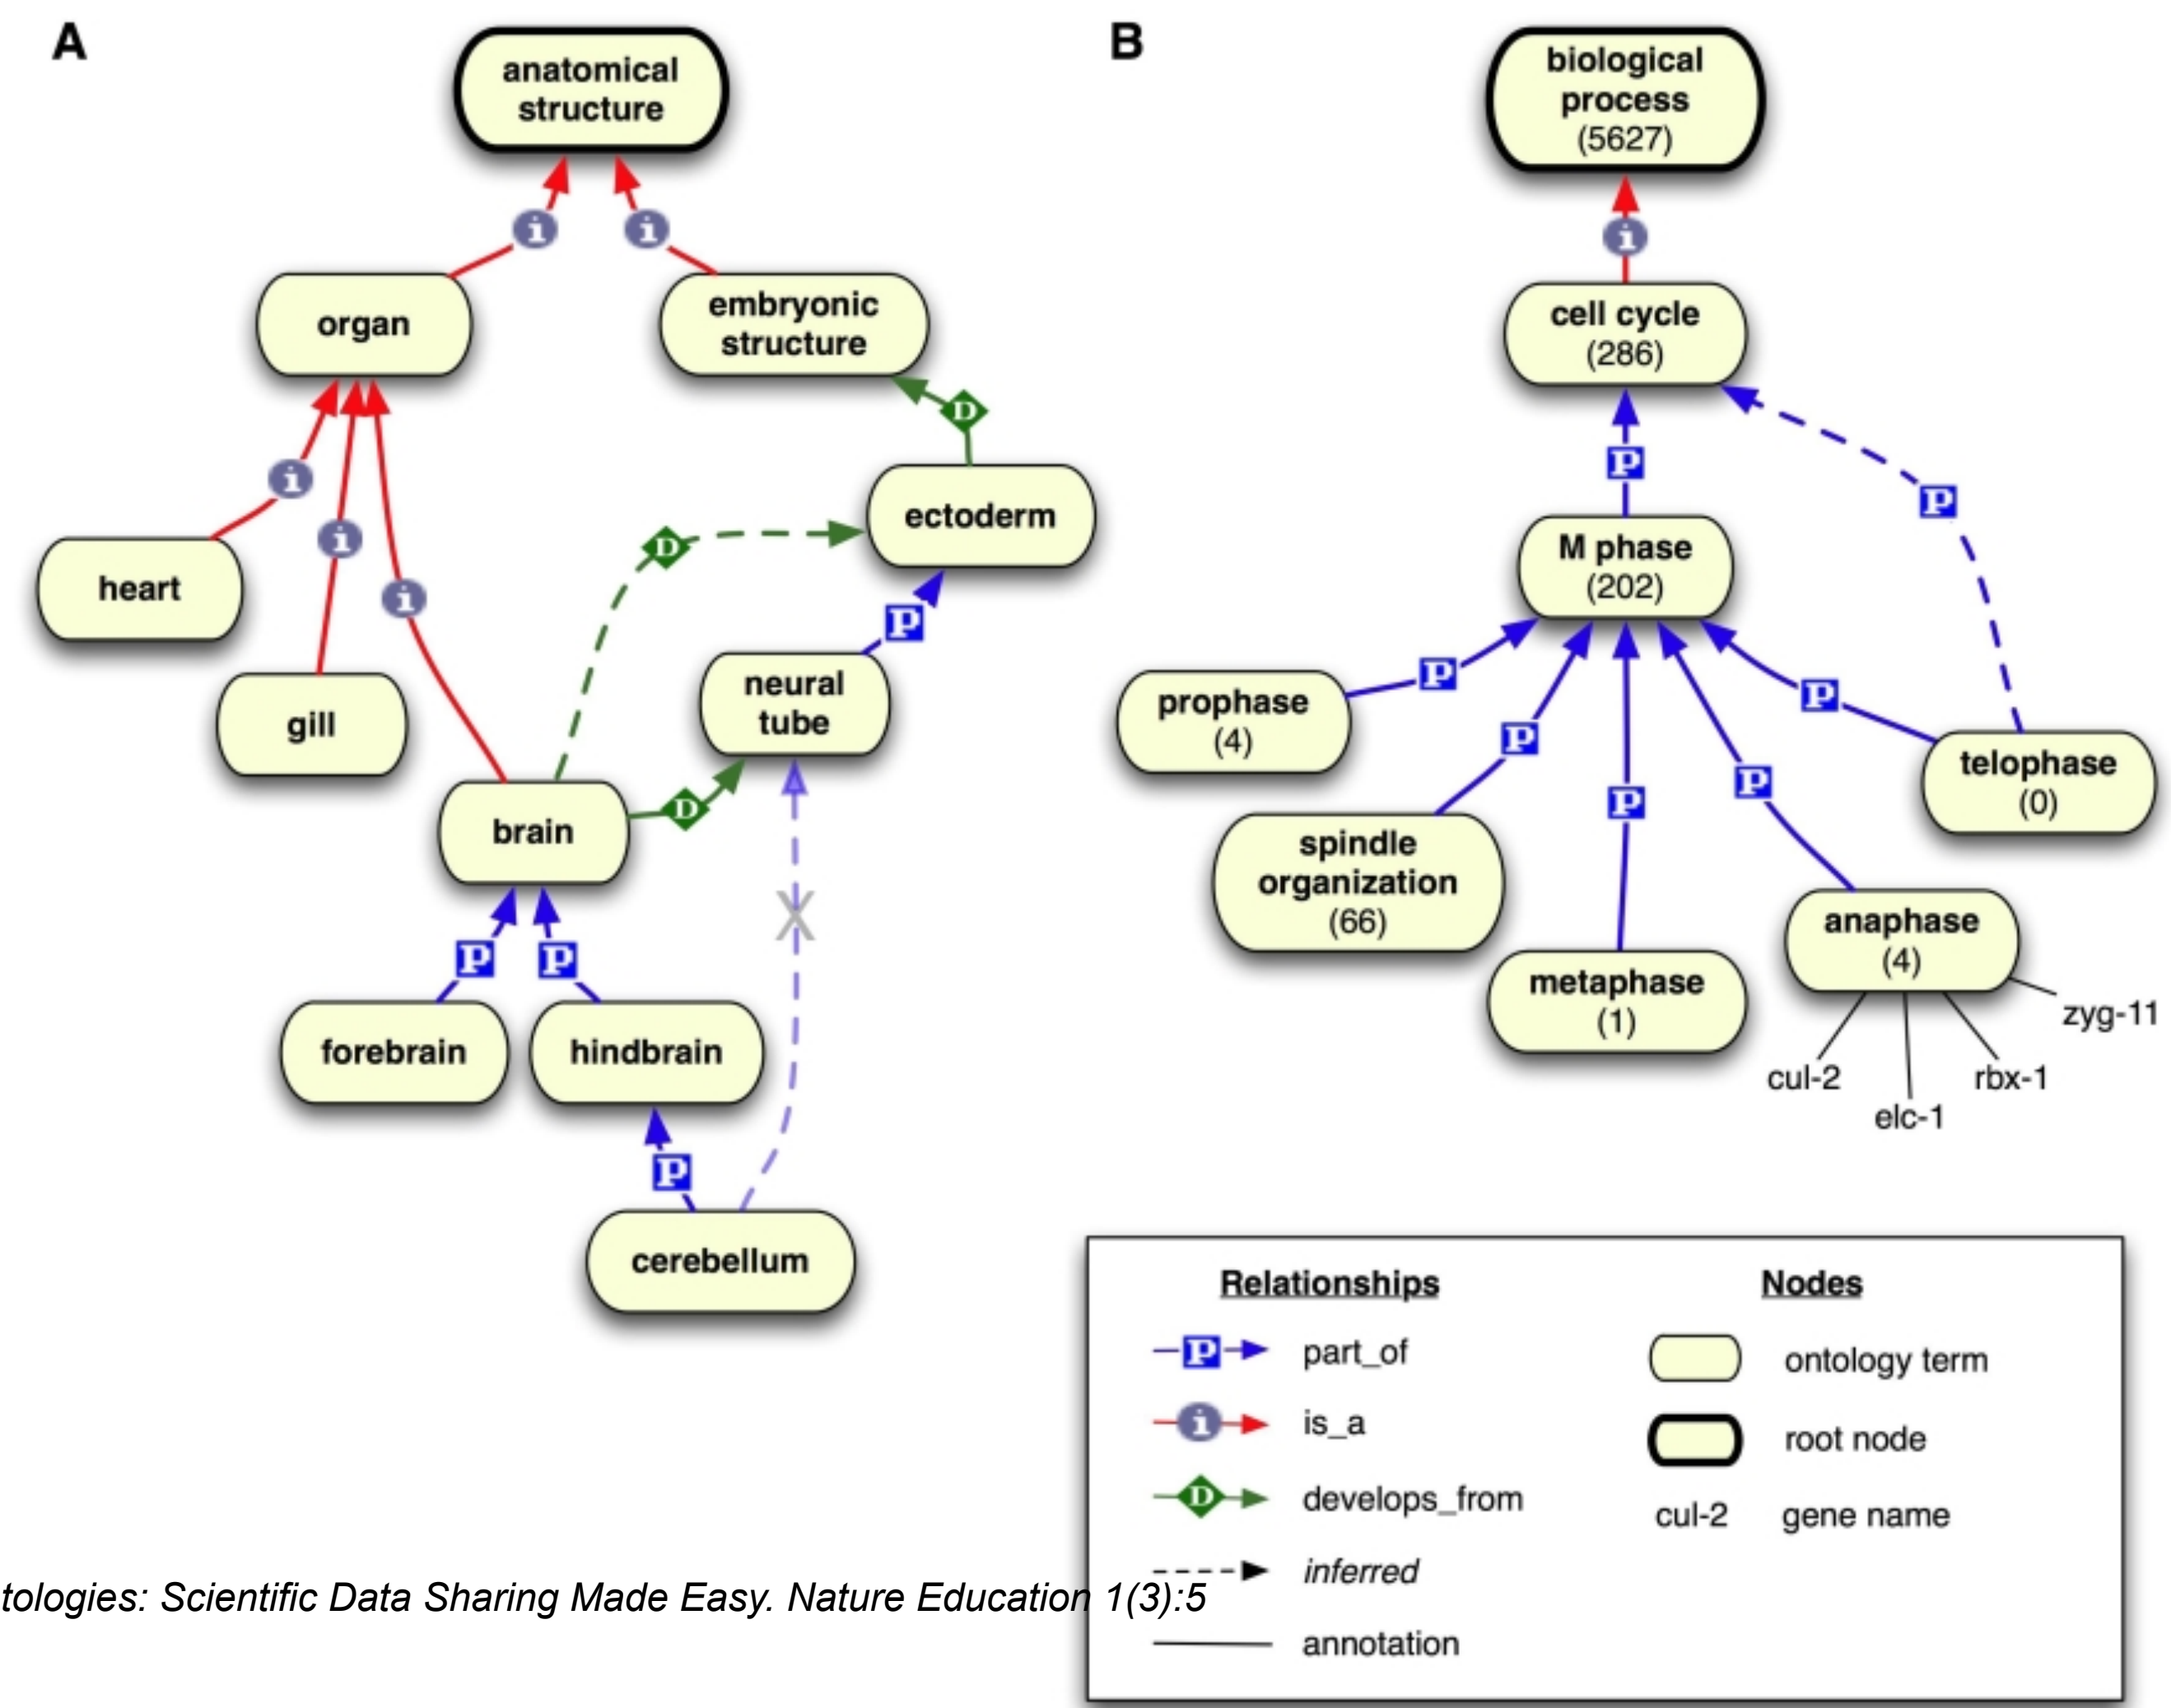
Elasticsearch Workshop Day

February 2016

Charlie Hull - Managing Director
Matt Pearce – Senior Developer

charlie@flax.co.uk
www.flax.co.uk/blog
+44 (0) 8700 118334
Twitter: @FlaxSearch

BioSolr & SPOt - Indexing ontologies



Washington, N. & Lewis, S. (2008) Ontologies: Scientific Data Sharing Made Easy. Nature Education 1(3):5

Indexing ontologies – the problem

- You have a collection of documents annotated with ontology references.
- You want to search both the documents and the associated ontology data.
 - This may include associated nodes – “has location”, “is part of”, etc.
- Faceting the ontology references would be nice!
 - (especially if the facets can be presented in a tree)

Approach 2 – final result

- We have developed a Solr UpdateProcessor to do this as part of the update chain.
- The user defines the ontology location (file, URL) and the field to reference.
- Aims for convention over configuration for remaining properties.
 - Field names and ontology annotation details all customisable.
- *A similar plugin for Elasticsearch has also been developed - we'll talk about this tomorrow!;*

ElasticSearch ontology annotator plugin

- Implemented as a mapper - a new field type.
- Adds new sub-fields to the annotation field.
- Configuration in field mapping, rather than index settings.

ElasticSearch ontology annotator plugin (2)

- Built on same ontology look-up code as the Solr version.
- (although code is currently duplicated, since we haven't published any mo
- Can use OWL/RDF files as base, or OLS API..

Differences from Solr implementation

- Configuration is (probably) simpler
- Adds sub-fields to the annotation field, rather than adding top-level fields.
- However, sub-fields not added to `_source`, need to be specified in query filter
- Installs using Elasticsearch plugin manager, rather than by copying JARs

Issues with this implementation

- Elasticsearch is constantly evolving
- Different code requirements for most versions of ES...
 - so we're aiming at a moving target, module will require constant updates
- Having to specify the new fields in your retrieval list is a pain.
- Elasticsearch community seem less supportive of plugin development than

Get Involved!

- Check out the github page: <https://github.com/flaxsearch/BioSolr>
- Suggest more use cases for what we've built!

Thank you for listening – any questions?

.co.uk/blog

700 118334

@flaxsearch